

CREATIVE PROGRAMMING AND COMPUTING

Lab 6:
Deep Learning

Luca Comanducci

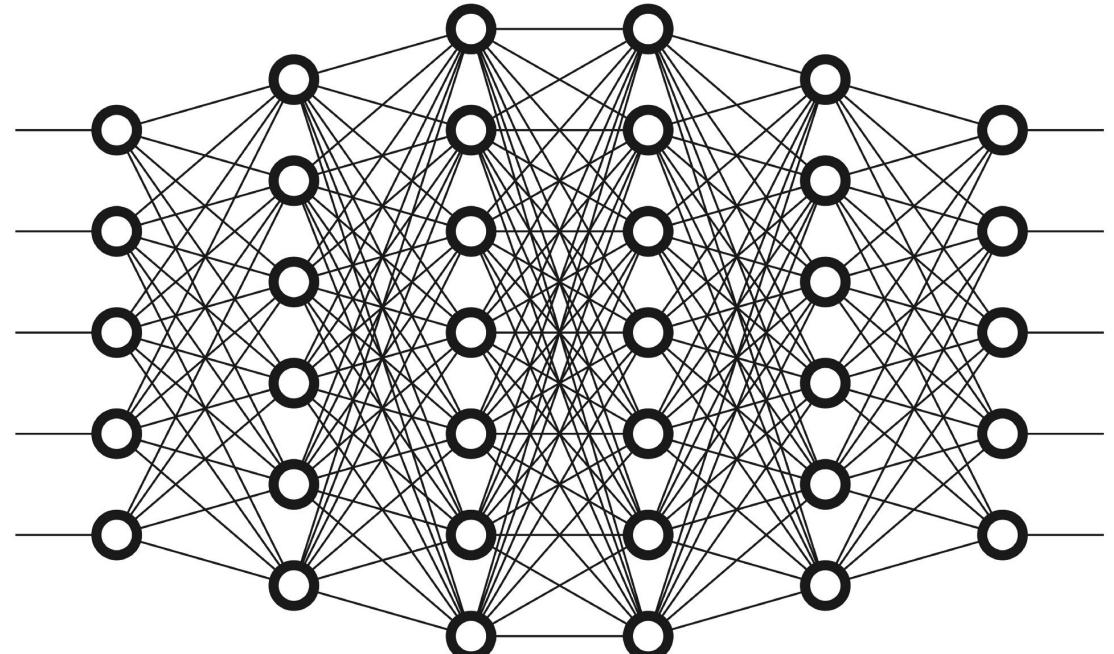
luca.comanducci@polimi.it

CONTENTS

- **Part I – Introduction to Keras and Tensorflow**
 - Basic Classification of Fashion images
- **Part II**
 - Neural Style Transfer
 - DeepDream

DEEP LEARNING

- Deep Learning is representation learning
- It is based on the small computation blocks called neural networks
- Neural Networks are organized into layers
- Deep learning methods automatically learn the representations (features) needed to fulfill the task considered (e.g. classification/regression)



DEEP LEARNING FRAMEWORKS

- *Tensorflow*



- *Keras*



- *PyTorch*



- *JAX*



- *Caffe*



- *Microsoft Cognitive Toolkit*

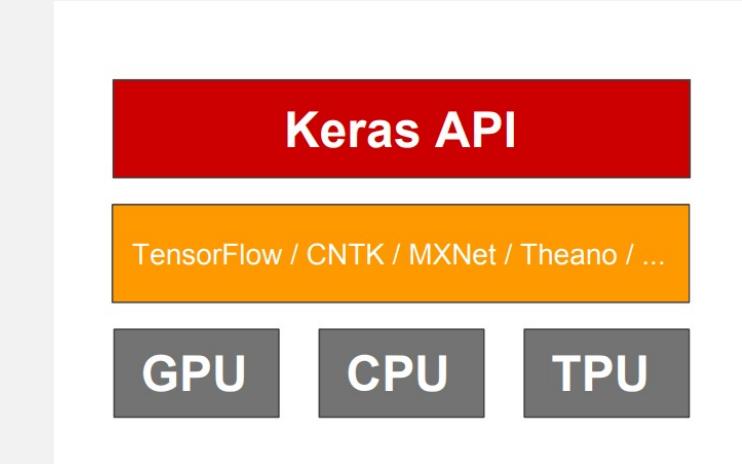


INTRODUCTION TO KERAS AND TENSORFLOW

INTRODUCTION TO KERAS AND TENSORFLOW

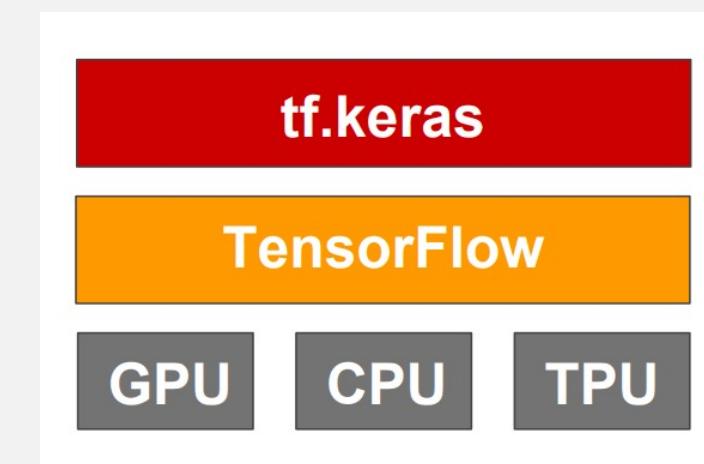
- **High-level neural networks API**

- Written in Python
- Runs on top of Tensorflow, Theano and CNTK
- User – friendly
- Runs both on CPU, GPU or TPU



- **Official high-level API of TensorFlow 2.0**

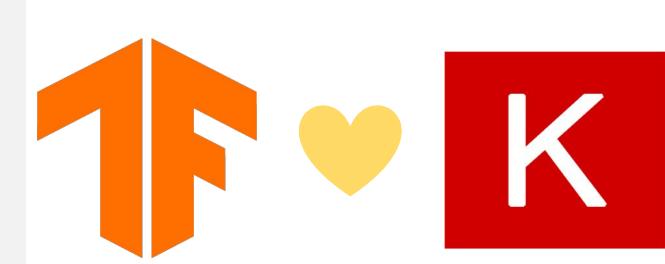
- tensorflow.keras module
- Better optimized for use with tensorflow
- Better integration with tensorflow specific features
 - tf.estimator
 - eager execution



INTRODUCTION TO KERAS AND TENSORFLOW

- Keras is built-in in tensorflow 2
- The right way to import is from tensorflow.keras

```
# !pip install tensorflow==2.0.0-beta1, then  
  
>>> from tensorflow.keras import layers # Right  
  
>>> from keras import layers # Oops  
  
Using TensorFlow backend. # You shouldn't see this
```



INTRODUCTION TO KERAS AND TENSORFLOW

- We will use several types of layers to build models, the most importants are
 - Fully connected layers `tf.keras.layers.Dense(units, activation = None, use_bias = True, kernel_initializer = 'glorot_uniform',....)`
 - Convolutional layers `tf.keras.layers.Conv2D(filters, kernel_size, strides=(1, 1), padding='valid', data_format=None,.....)`
 - Flatten layers `tf.keras.layers.Flatten()`
 - Activation functions can be e.g. "relu", "sigmoid", "tanh"
- Three model building styles
 - Sequential
 - Functional
 - Subclassing

INTRODUCTION TO KERAS AND TENSORFLOW

- **Sequential-Model**

- Extremely simple
- Good only for one-input one-output stacks

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(M,N),name='flatten'),
    tf.keras.layers.Dense(512, activation='relu', name='fc1'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax',name='output')
])
```

INTRODUCTION TO KERAS AND TENSORFLOW

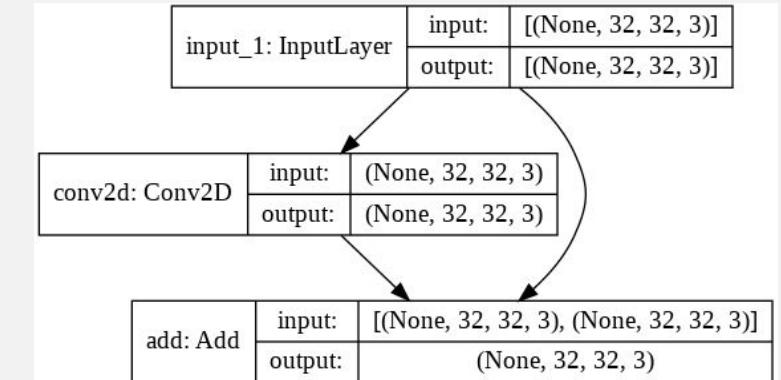
- **Functional models**

- Multi-input multi-output arbitrary topologies

```
inputs = keras.Input(shape=(32, 32, 3))
```

```
y = layers.Conv2D(3, (3, 3), activation='relu', padding='same')(inputs)  
outputs = layers.add([inputs, y])
```

```
model = keras.Model(inputs, outputs)
```



INTRODUCTION TO KERAS AND TENSORFLOW

- **Subclassed Models**

- Multi-input multi-output arbitrary topologies

```
class MyModel(tf.keras.Model):  
    def __init__(self, num_classes=10):  
        super(MyModel, self).__init__(name='my_model')  
        self.dense_1 = layers.Dense(32, activation='relu')  
        self.dense_2 = layers.Dense(num_classes, activation='sigmoid')  
  
    def call(self, inputs):  
        # Define your forward pass here  
        x = self.dense_1(inputs)  
        return self.dense_2(x)
```

INTRODUCTION TO KERAS AND TENSORFLOW

- **Compile**

- After a model has been defined, it must be compiled, specifying the optimizer, loss, metrics and eventual callbacks

```
model.compile(  
    optimizer='adam',  
    loss='sparse_categorical_crossentropy',  
    metrics=['accuracy'])
```

INTRODUCTION TO KERAS AND TENSORFLOW

- **Training**

- Once a model is written and compiled it can be trained in two main different ways
 - **Built-in training loop**

```
model.fit(x_train, y_train, epochs=5)
```

- **Custom-loop**

```
model = MyModel()  
with tf.GradientTape() as tape:  
    logits = model(images)  
    loss_value = loss(logits, labels)  
  
grads = tape.gradient(loss_value, model.trainable_variables)  
optimizer.apply_gradients(zip(grads, model.trainable_variables))
```

INTRODUCTION TO KERAS AND TENSORFLOW

- Example of full model:

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

HOW TO RUN THESE EXERCISES?

- You will find some jupyter notebooks on the github of the course
- Jupyter is a web-based interactive development environment for Jupyter notebooks



- You can use Google Colaboratory to run them



- Colaboratory is a free Jupyter notebook environment that requires no setup and runs entirely in the cloud and allows you to access powerful computing resources.

INTRODUCTION TO KERAS

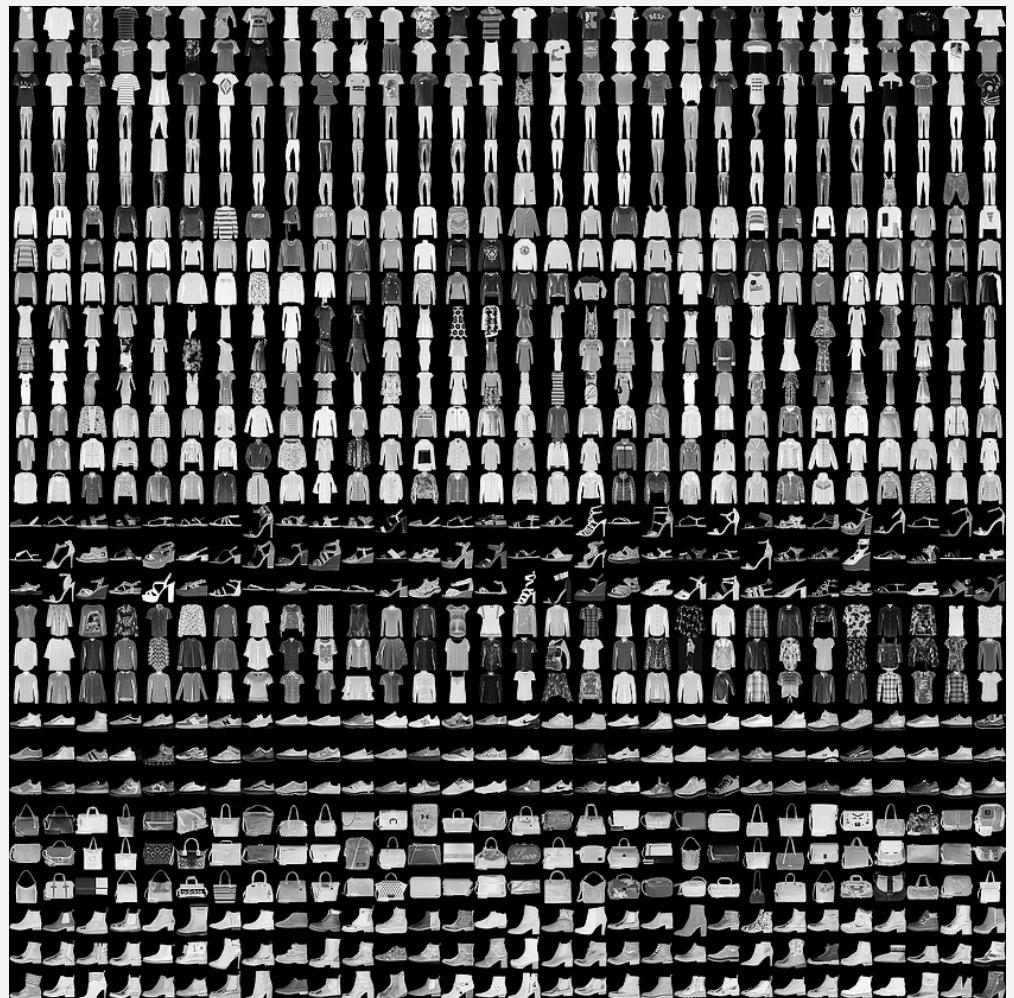
- Basic Classification of clothing images
- You will find the code in the file:
“Ex.1 - Introduction to Keras - Classification of images of clothing.ipynb”

CLASSIFICATION OF CLOTHING IMAGES

- We will consider the MNIST fashion dataset
 - 70,000 grayscale images in 10 categories
 - Replacement to MNIST dataset (benchmark for models)
 - MNIST dataset contains images of handwritten digits (0, 1, 2, etc.) in a format identical to that of the articles of clothing



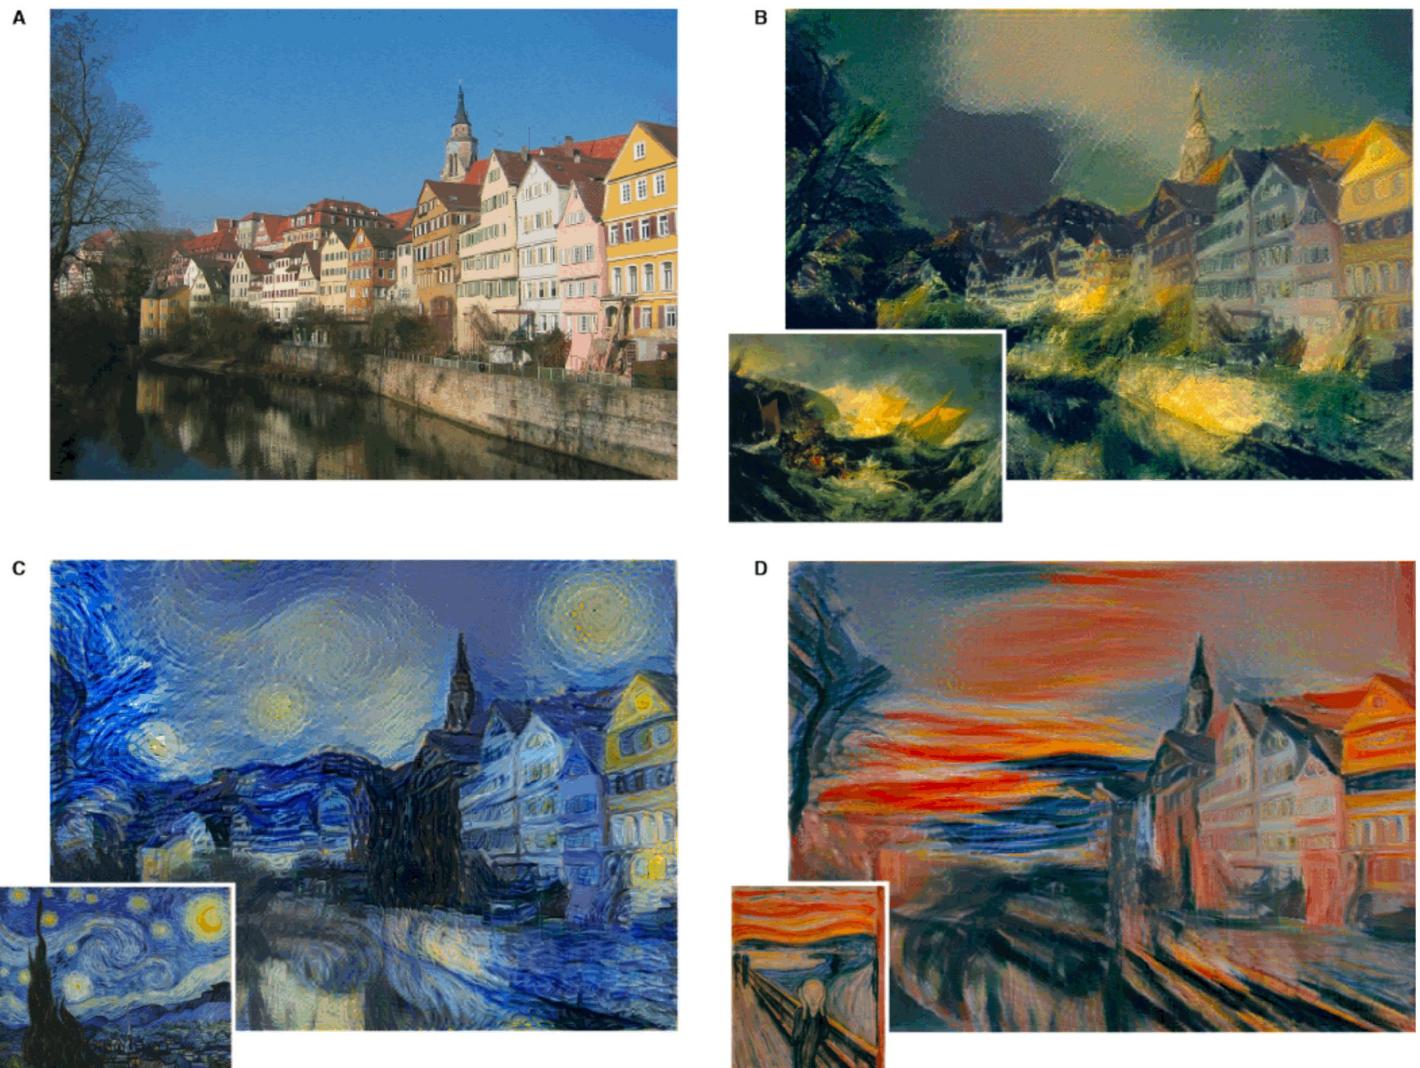
- You will learn how to train a network based only on feedforward NN, in order to classify the images



NEURAL STYLE TRANSFER

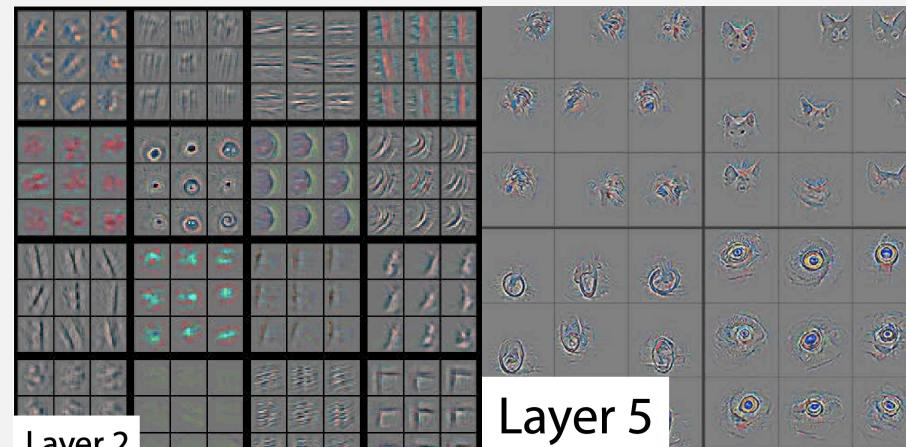
NEURAL STYLE TRANSFER

- Deep Neural Network that creates artistic images of high perceptual quality
- Separates and recombines the content and style of an image
- representations of content and style in the Convolutional Neural Networks (CNNs) are separable.

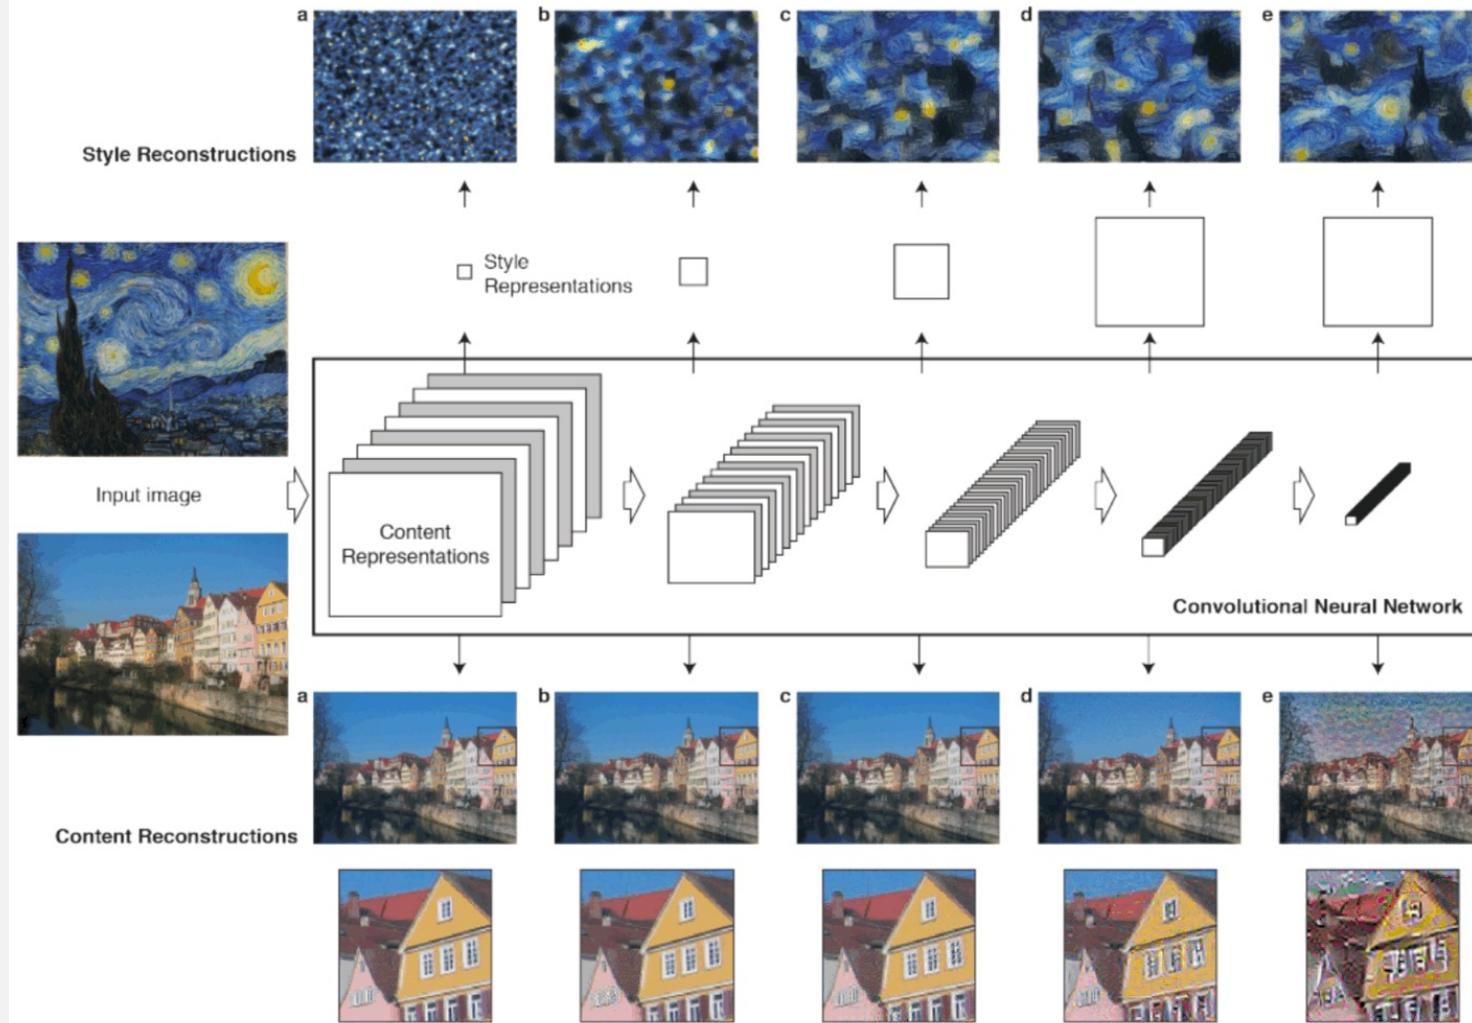


NEURAL STYLE TRANSFER

- CNNs consist of layers of small computational units that process information hierarchically in a feed-forward manner
- Each layer extract a certain feature-maps from the image
- CNNs trained for object recognition
 - Lower layers: learn low-level spatial feature
 - Higher layers: capture high-level content in terms of objects and their arrangement (**Content**)
 - Extracting the correlations between the different filter responses over the spatial extent of the feature maps we obtain a stationary, multi-scale representation of the input image, which captures its texture information but not the global arrangement (**Style**)

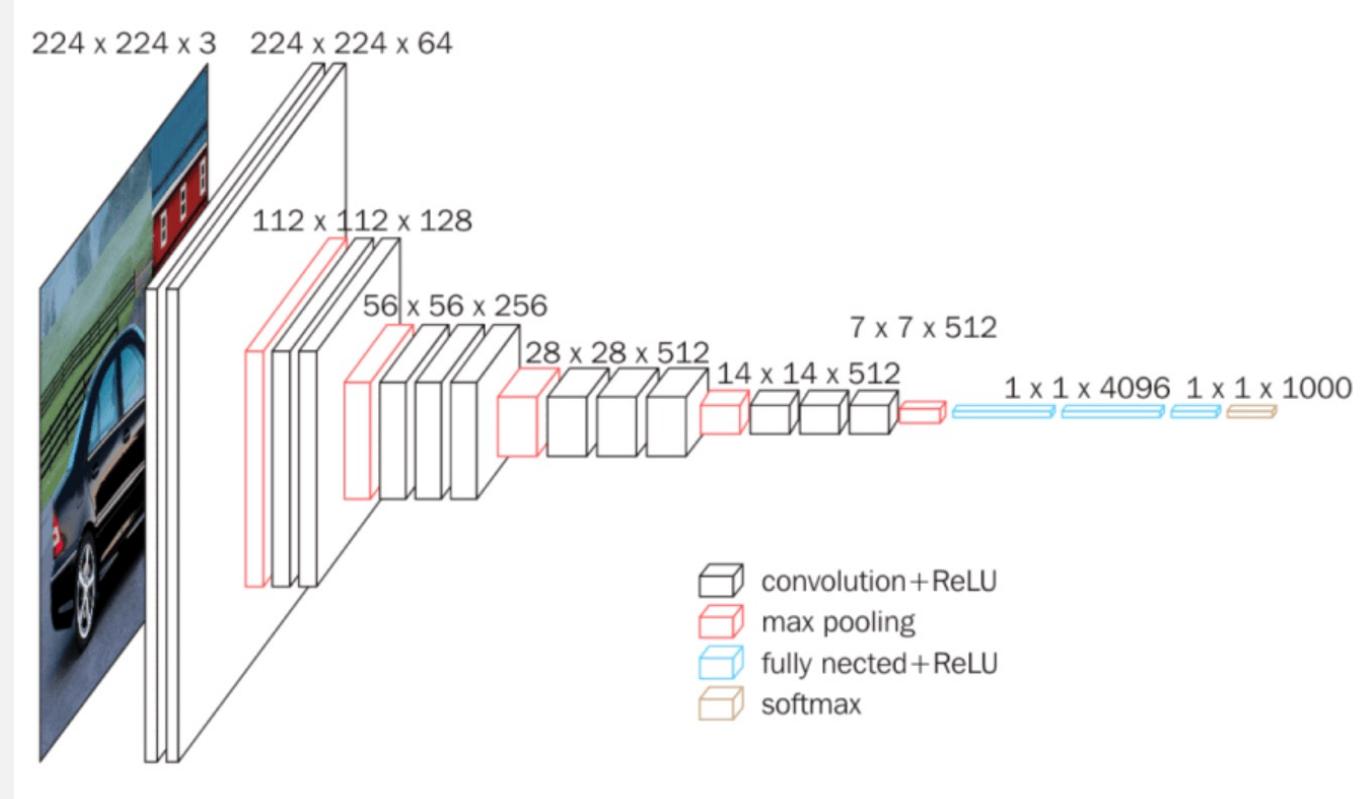


NEURAL STYLE TRANSFER



NEURAL STYLE TRANSFER

- images are synthesised by finding an image that simultaneously matches the content representation of the photograph and the style representation of the respective piece of art
- Features are extracted using the pre-trained classification model
VGG
- Achieved 92.7 % accuracy in 2014 on the ImageNet dataset (1000 classes, 14 million images)



NEURAL STYLE TRANSFER

- each layer in the network defines a non-linear filter bank whose complexity increases with the position of the layer in the network
- A layer with N_l filters has N_l feature maps, each of size $M_l = \text{height} * \text{width}$
- Response in layer l can be stored in a matrix

where F_{ij}^l contains the activation of the i-th filter at position j in layer l

$$F^l \in \mathcal{R}^{N_l \times M_l}$$

NEURAL STYLE TRANSFER

- ***Content loss:***

- We use gradient descent to recover the content of the image
 - $P \rightarrow$ original image $P \rightarrow$ feature-maps
 - $x \rightarrow$ generated image $X \rightarrow$ feature-maps
- We define the loss as the squared error between the two representations

$$\mathcal{L}_{content}(\vec{p}, \vec{x}, l) = \frac{1}{2} \sum_{i,j} (F_{ij}^l - P_{ij}^l)^2$$

NEURAL STYLE TRANSFER

- **Style loss:**

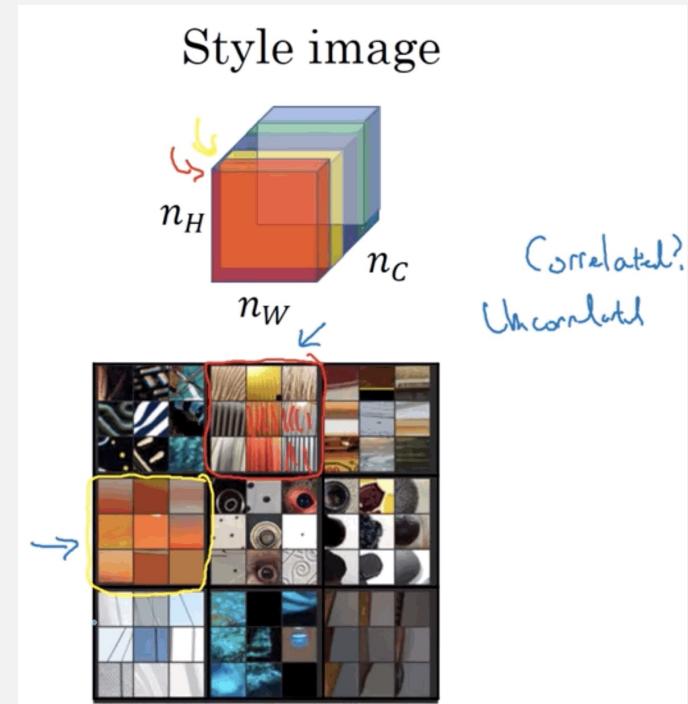
- Compute the correlation between the different filter responses, where the expectation is taken over the spatial extent of the input image

- We use the Gram matrix

$$G^l = \sum_k F_{ik}^l F_{jk}^l$$

- Through gradient descent we minimize the mean-squared distances between the Gram matrix of the original image A and of the generated one B

$$\mathcal{L}_{style}(\vec{a}, \vec{x}) = \frac{1}{4N_l^2 M_l^2} \sum_{l=0}^L \sum_{i,j} (G_{ij}^l - A_{ij}^l)^2$$

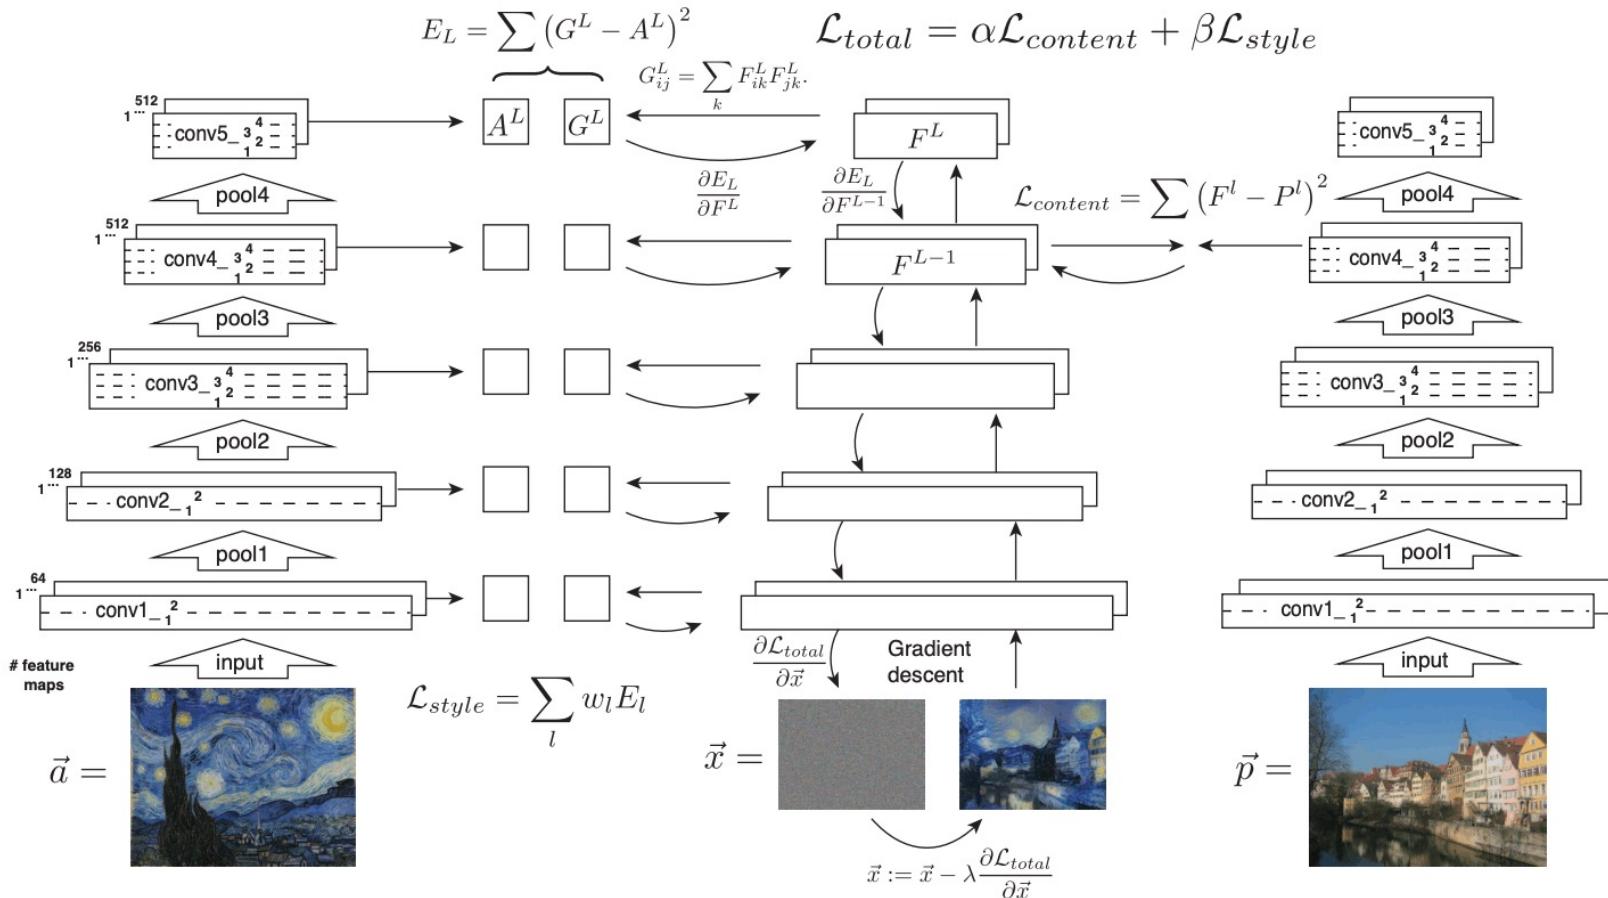


NEURAL STYLE TRANSFER

- ***Full loss:***
 - Compute the correlation between the different filter responses, where the expectation is taken over the spatial extent of the input image

$$\mathcal{L}_{total}(\vec{p}, \vec{a}, \vec{x}) = \alpha \mathcal{L}(\vec{p}, \vec{x}) + \beta \mathcal{L}_{style}(\vec{a}, \vec{x})$$

NEURAL STYLE TRANSFER

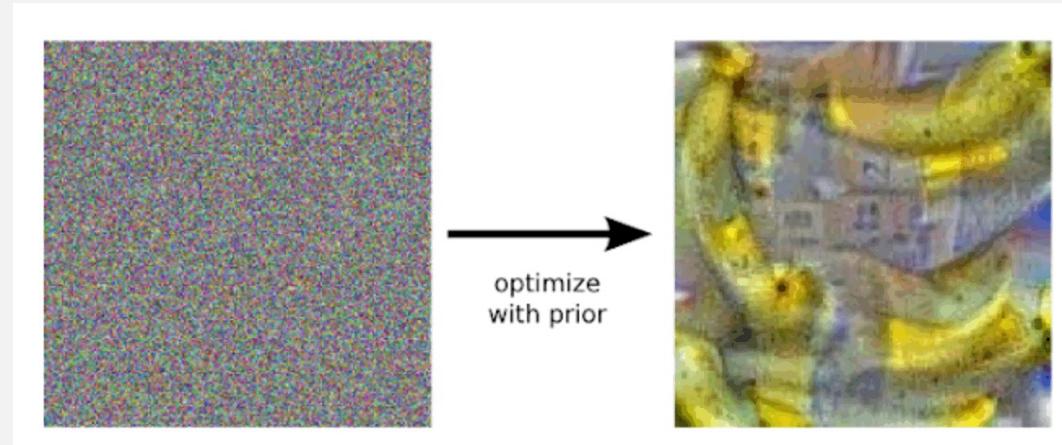


DEEPDREAM

- You will find the code in the file:
“*DeepDream.ipynb*”

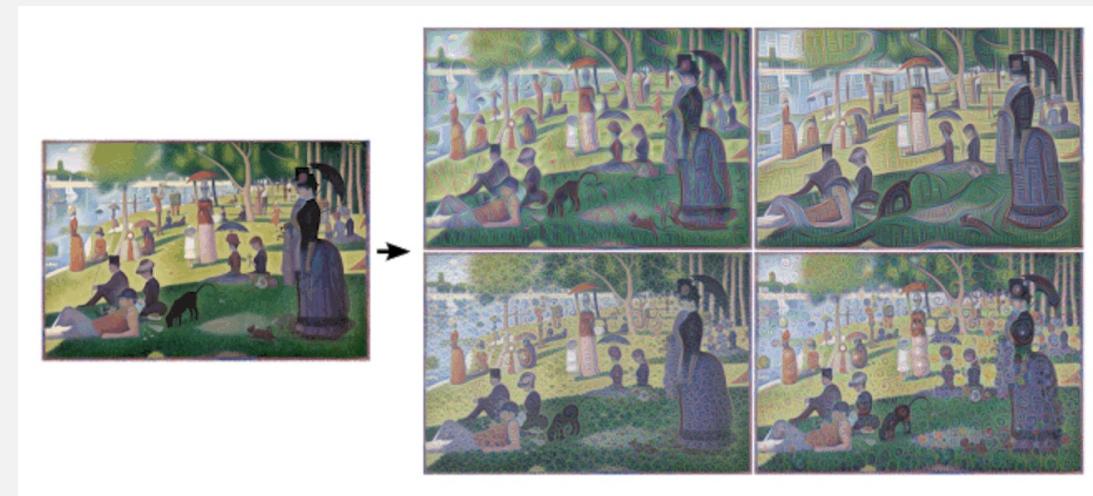
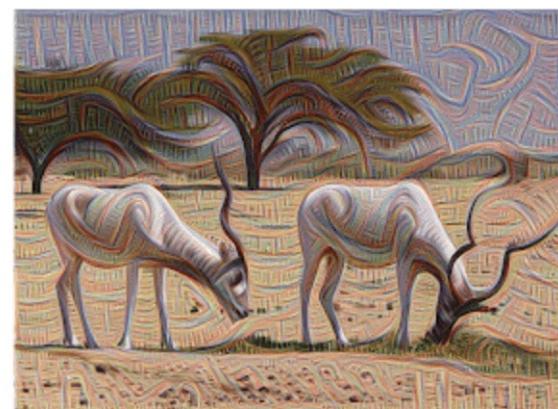
DEEPDREAM

- One of the challenges of neural networks is understanding what exactly goes on at each layer.
- In classification we know that after training, each layer progressively extracts higher and higher-level features of the image, until the final layer essentially makes a decision on what the image shows
- **Idea:**
 - turn the network upside down and ask it to enhance an input image in such a way as to elicit a particular interpretation
 - One way to visualize what goes on is to turn the network upside down and ask it to enhance an input image in such a way as to elicit a particular interpretation



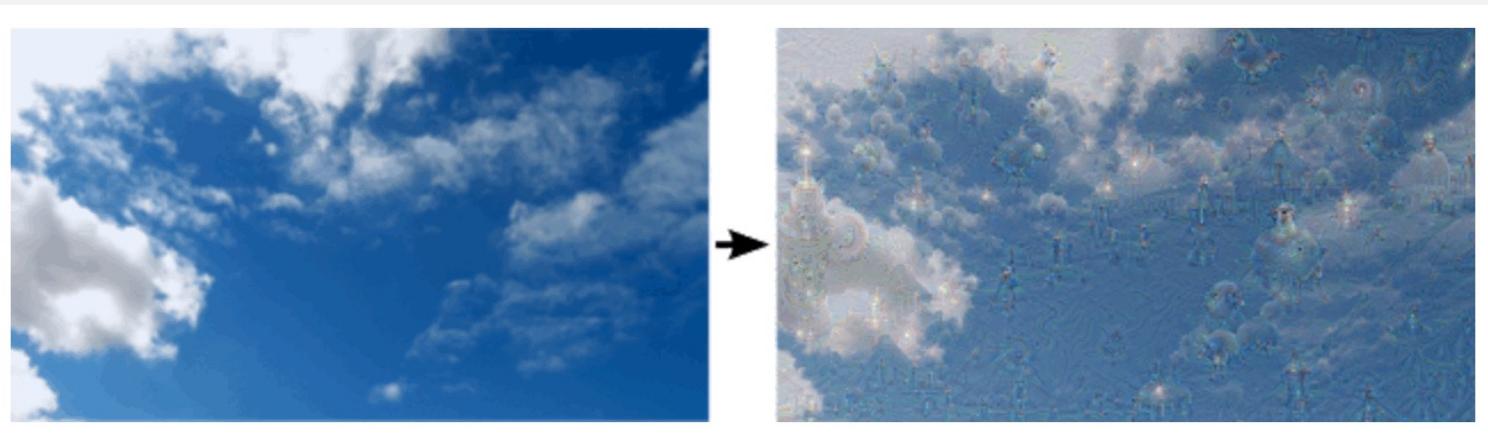
DEEPDREAM

- Instead of exactly prescribing which feature we want the network to amplify, we can also let the network make that decision.
- We feed the network an arbitrary image or photo and let the network analyze the picture. We then **pick a layer and ask the network to enhance whatever it detected**.
- Each layer of the network deals with features at a different level of abstraction
 - **Lower layers** -> strokes or simple ornament-like patterns



DEEPDREAM

- **Higher layers** -> identify more sophisticated features in images
 - complex features or even whole objects tend to emerge



"Admiral Dog!"



"The Pig-Snail"



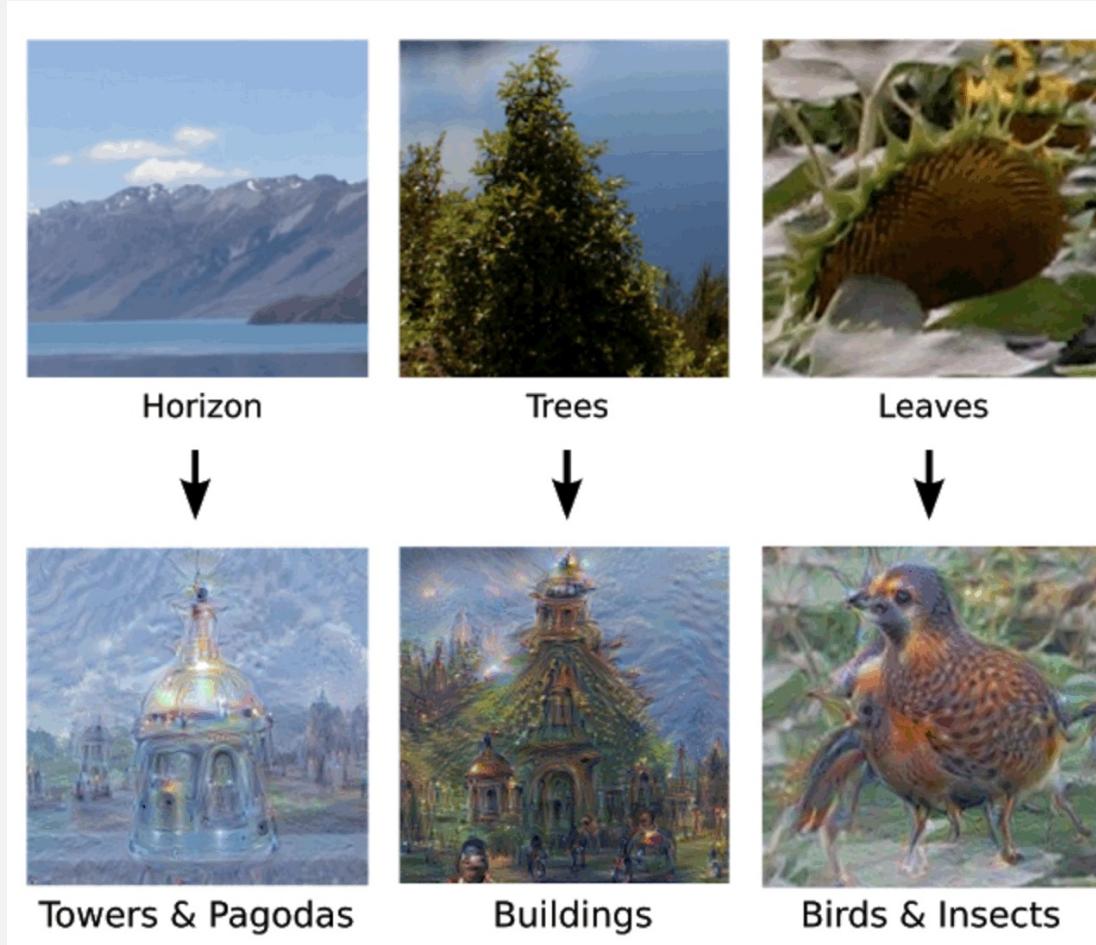
"The Camel-Bird"



"The Dog-Fish"

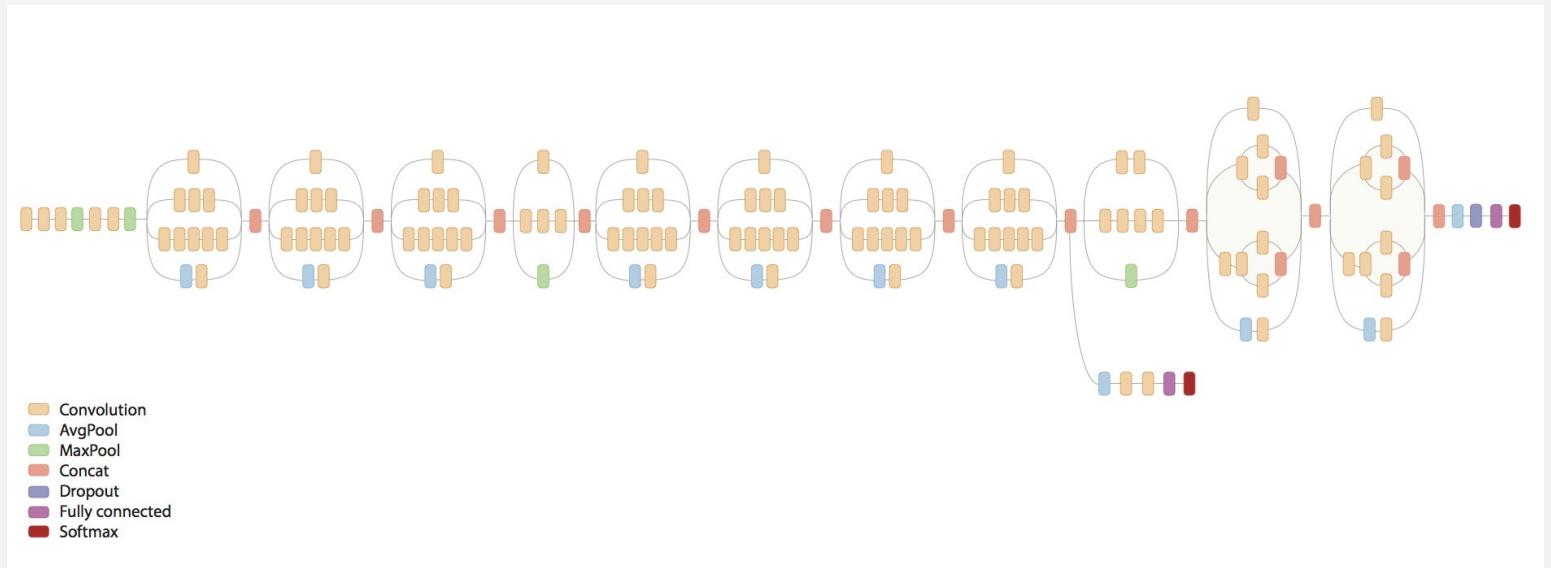
DEEPDREAM

- We can apply this technique to any kind of image. Results vary depending on the kind of image considered



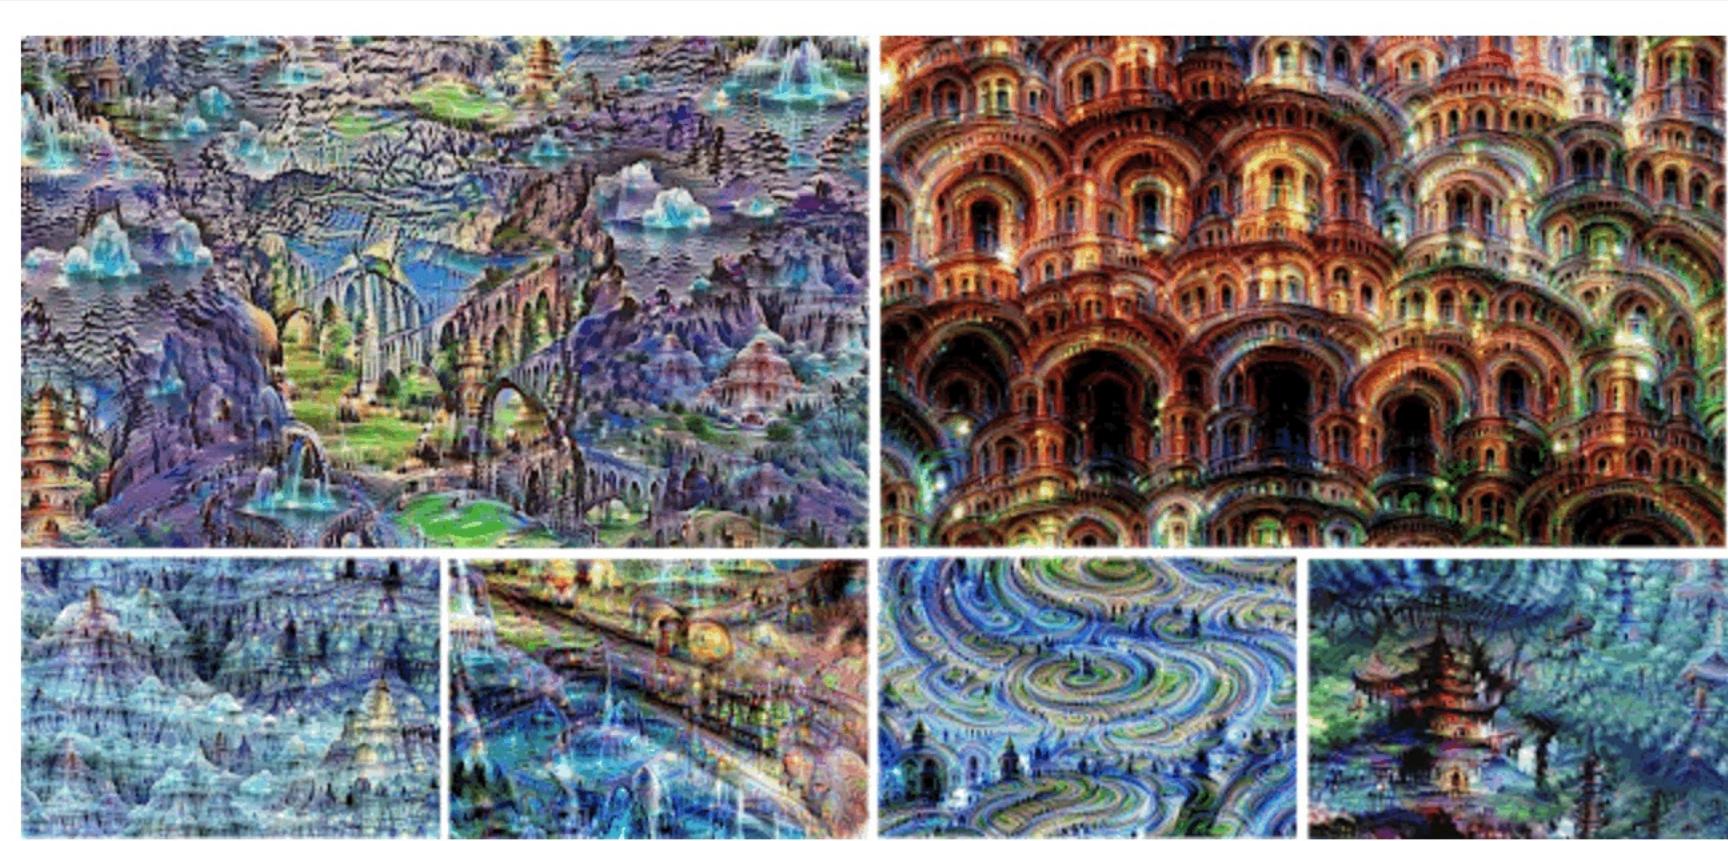
DEEPDREAM

- To wrap up,
 - DeepDream Idea:
 - Feed a pre-trained classification network with images
 - Maximize the activations of chosen layers using **gradient ascent**
 - the technique is also called **Inceptionism** referring to the pre-trained architecture used
 - Quite large architecture



DEEPDREAM

- If we repeat the algorithm endlessly on its own outputs and apply zooming, we get an endless stream of new impressions.



REFERENCES

- **Deep Learning**
 - <http://www.deeplearningbook.org/>
- **Style Transfer**
 - Gatys, Leon A., Alexander S. Ecker, and Matthias Bethge, [A neural algorithm of artistic style](#)
 - Simonyan, Karen, and Andrew Zisserman, [Very deep convolutional networks for large-scale image recognition.](#)
 - [Tensorflow tutorial](#)
 - [Fast style transfer on webcam](#), Gene Kogan
- **DeepDream**
 - [Inceptionism: Going Deeper into Neural Networks](#)
 - Szegedy, Christian, et al. [Going deeper with convolutions](#)
 - [Tensorflow tutorial](#)