

CREATIVE PROGRAMMING AND COMPUTING

Lab 7:

Time-aware generative models and Generative Adversarial Networks (GANs)

Luca Comanducci

luca.comanducci@polimi.it

CONTENTS

- **Part I**
 - Music generation using Recurrent Neural Networks (RNNs)
- **Part II**
 - Image-to-Image Translation with Conditional Adversarial Networks (Pix2Pix)

PART I: MUSIC GENERATION WITH RECURRENT NEURAL NETWORKS

- You can find the code in the file:
"Ex1 Music Generation with RNNs.ipynb"

EXERCISE I: MUSIC GENERATION WITH RECURRENT NEURAL NETWORKS

- *Objective:* apply Recurrent Neural Networks (RNNs) for music generation
- *How:* we will train a model to learn patterns contained in raw sheet music in ABC notation, treating it as a sequence of characters.
- We will consider a dataset made of Irish Folk songs

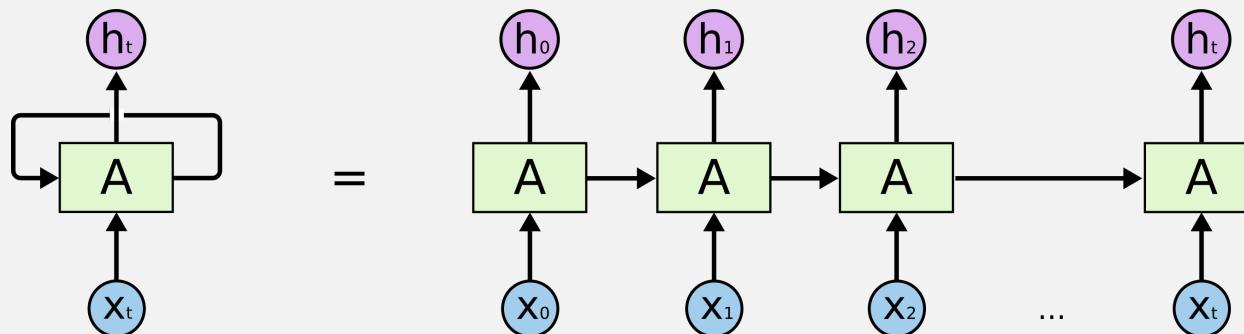
X: 0
T: My Tune
M: 4/4
L: 1/4
K: C
C,D,E,F, | G,A,B,C | D E F G | A B C d | e f g a | b c 'd 'e' | f 'g 'a 'b' |

My Tune



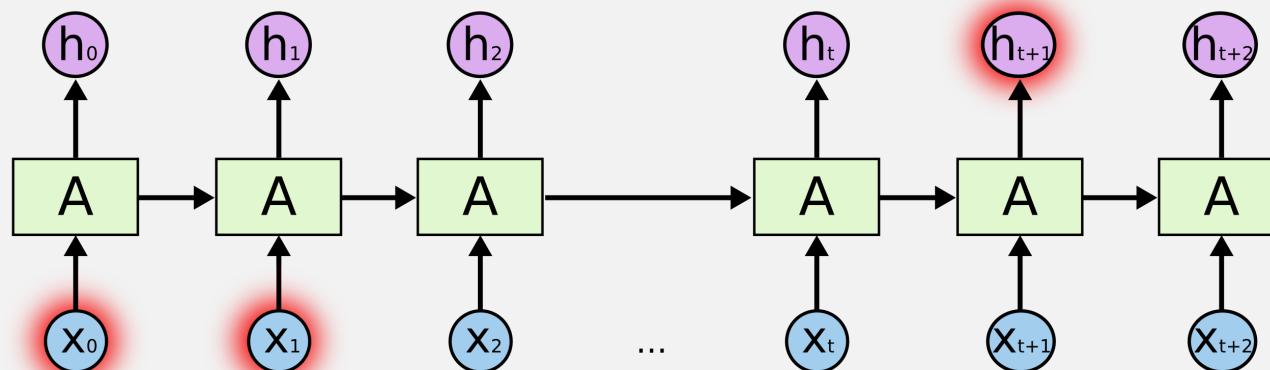
EXERCISE I: MUSIC GENERATION WITH RECURRENT NEURAL NETWORKS

- We will use Long Short Term Memory network Cells in order to model long term dependencies
- RNNs are good at connecting previous information to the current task, but what about long-term dependencies?



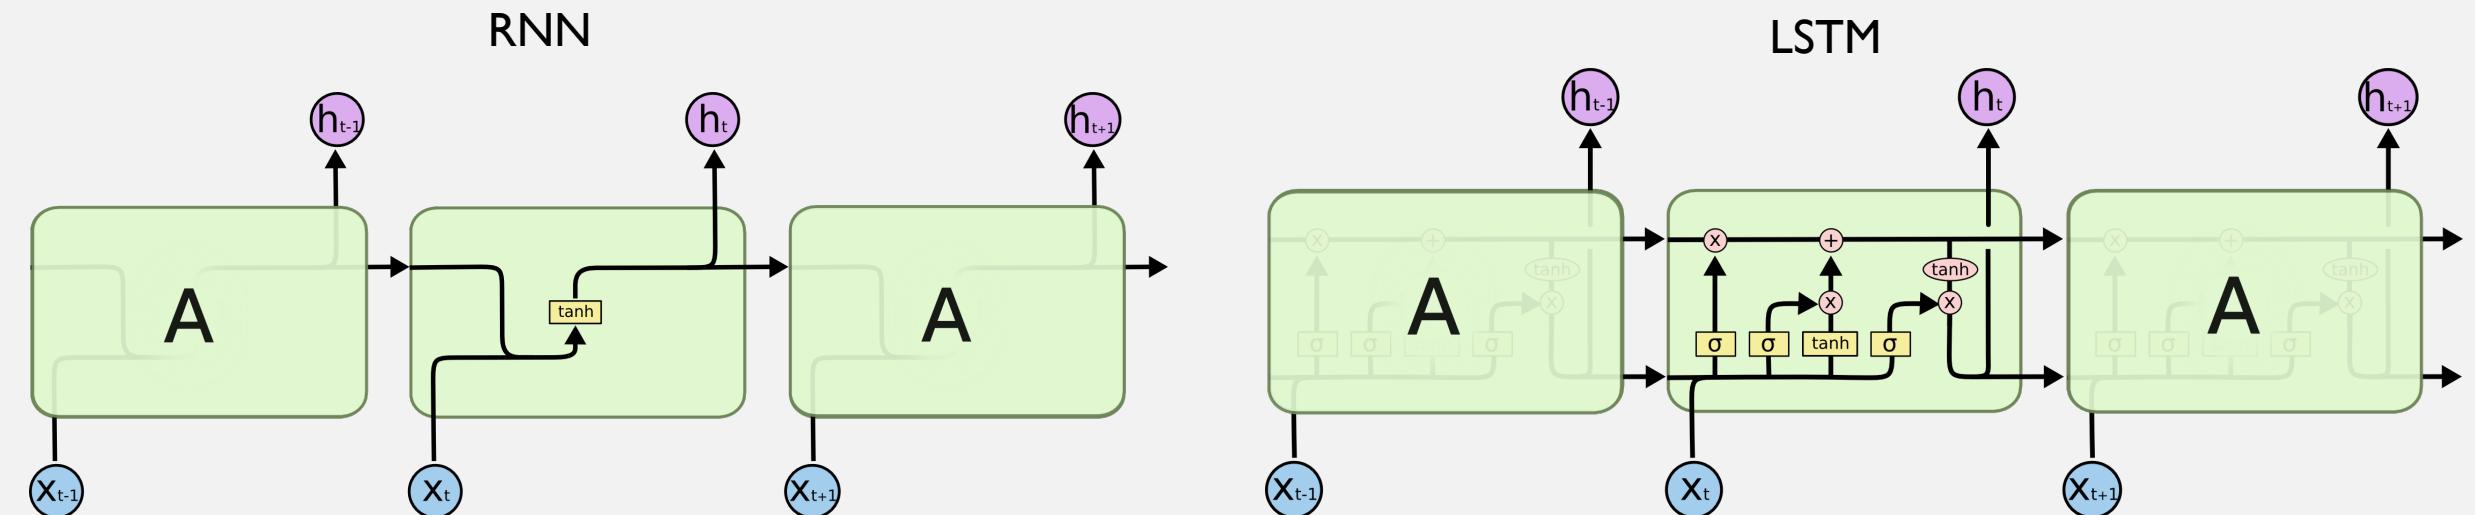
EXERCISE I: MUSIC GENERATION WITH RECURRENT NEURAL NETWORKS

- E.g. imagine you want to predict the last word "French" in a long text: "I grew up in France, ... I speak fluent French"
- Recent info suggests: the word is probably the name of a language
- Long-term info suggests: French
- As the gap grows bigger RNNs struggle to connect the information



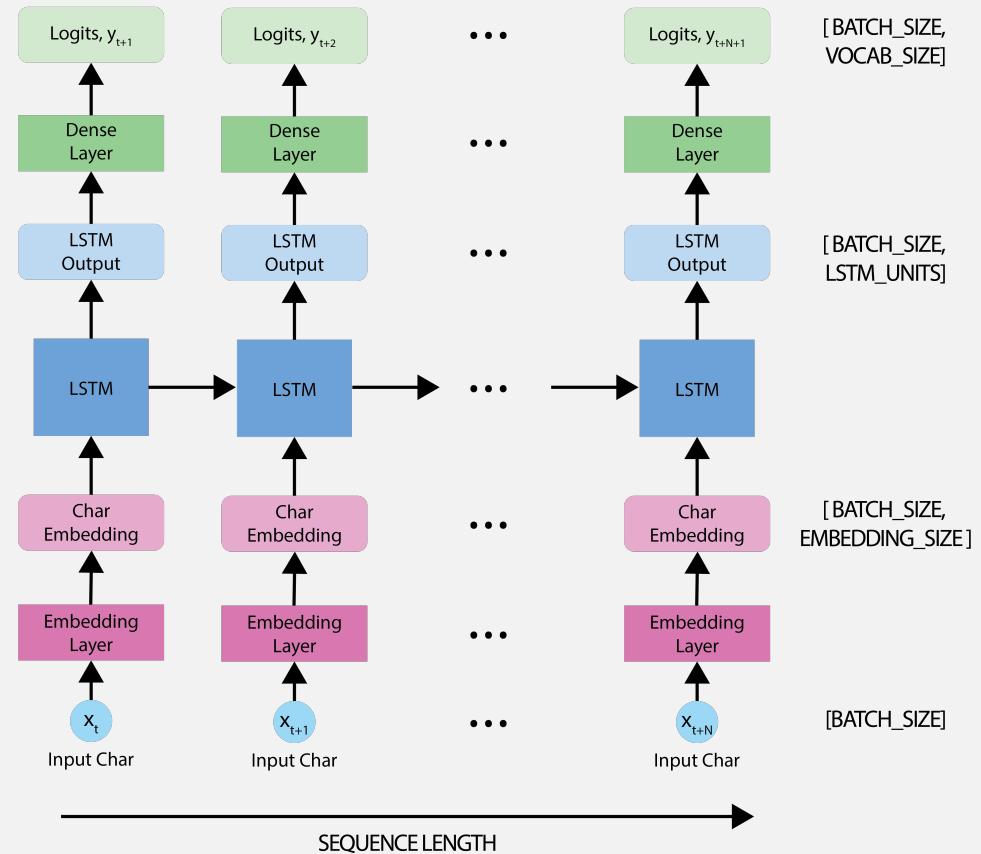
EXERCISE I: MUSIC GENERATION WITH RECURRENT NEURAL NETWORKS

- This is why Long Short Term Memory (LSTM) networks were introduced
- The fundamental idea is that LSTM are able to select which info must pass through at each time instant



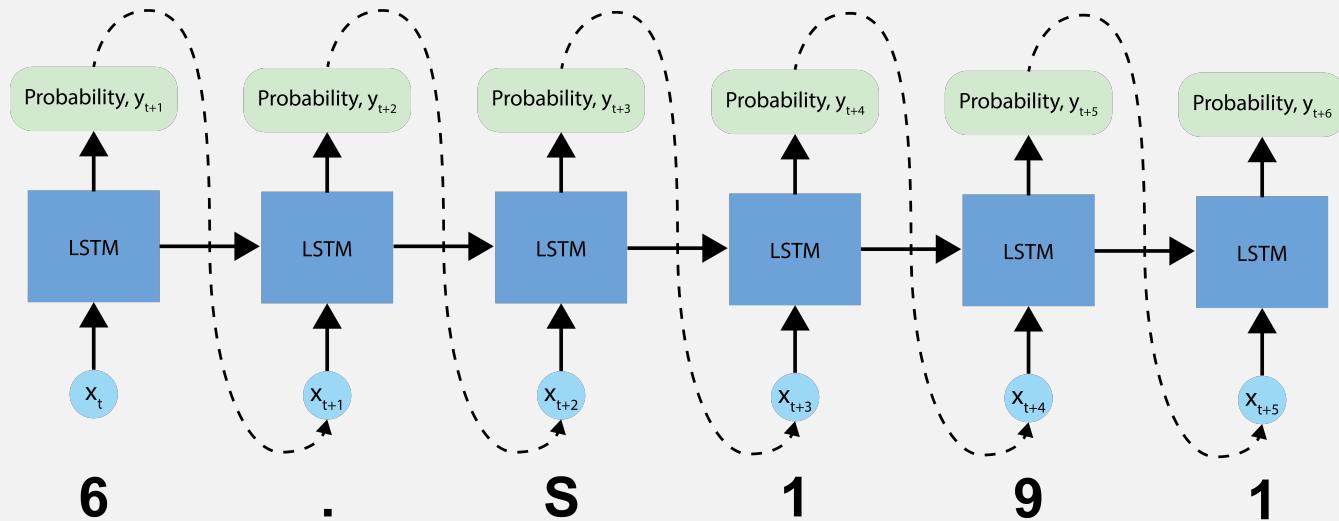
EXERCISE I: MUSIC GENERATION WITH RECURRENT NEURAL NETWORKS

- How does the model work?
 - One-to-one mapping between ABC notation and numbers
 - Network gets a sequence of characters and outputs the probability of each character in being the next one
 - This means that we are generating an output distribution corresponding to a categorical distribution (Generalized Bernoulli Distribution), from which we can sample the next character (i.e. note or other symbol in ABC notation)



EXERCISE I: MUSIC GENERATION WITH RECURRENT NEURAL NETWORKS

- How does the prediction procedure work?
 - Initialize a “seed” start string to obtain probability distribution over next predicted character.
 - Sample from the multinomial distribution to calculate next predicted character, which will be next input to model.
 - At each time step updated RNN state is fed back into the model, learning sequence dependencies.

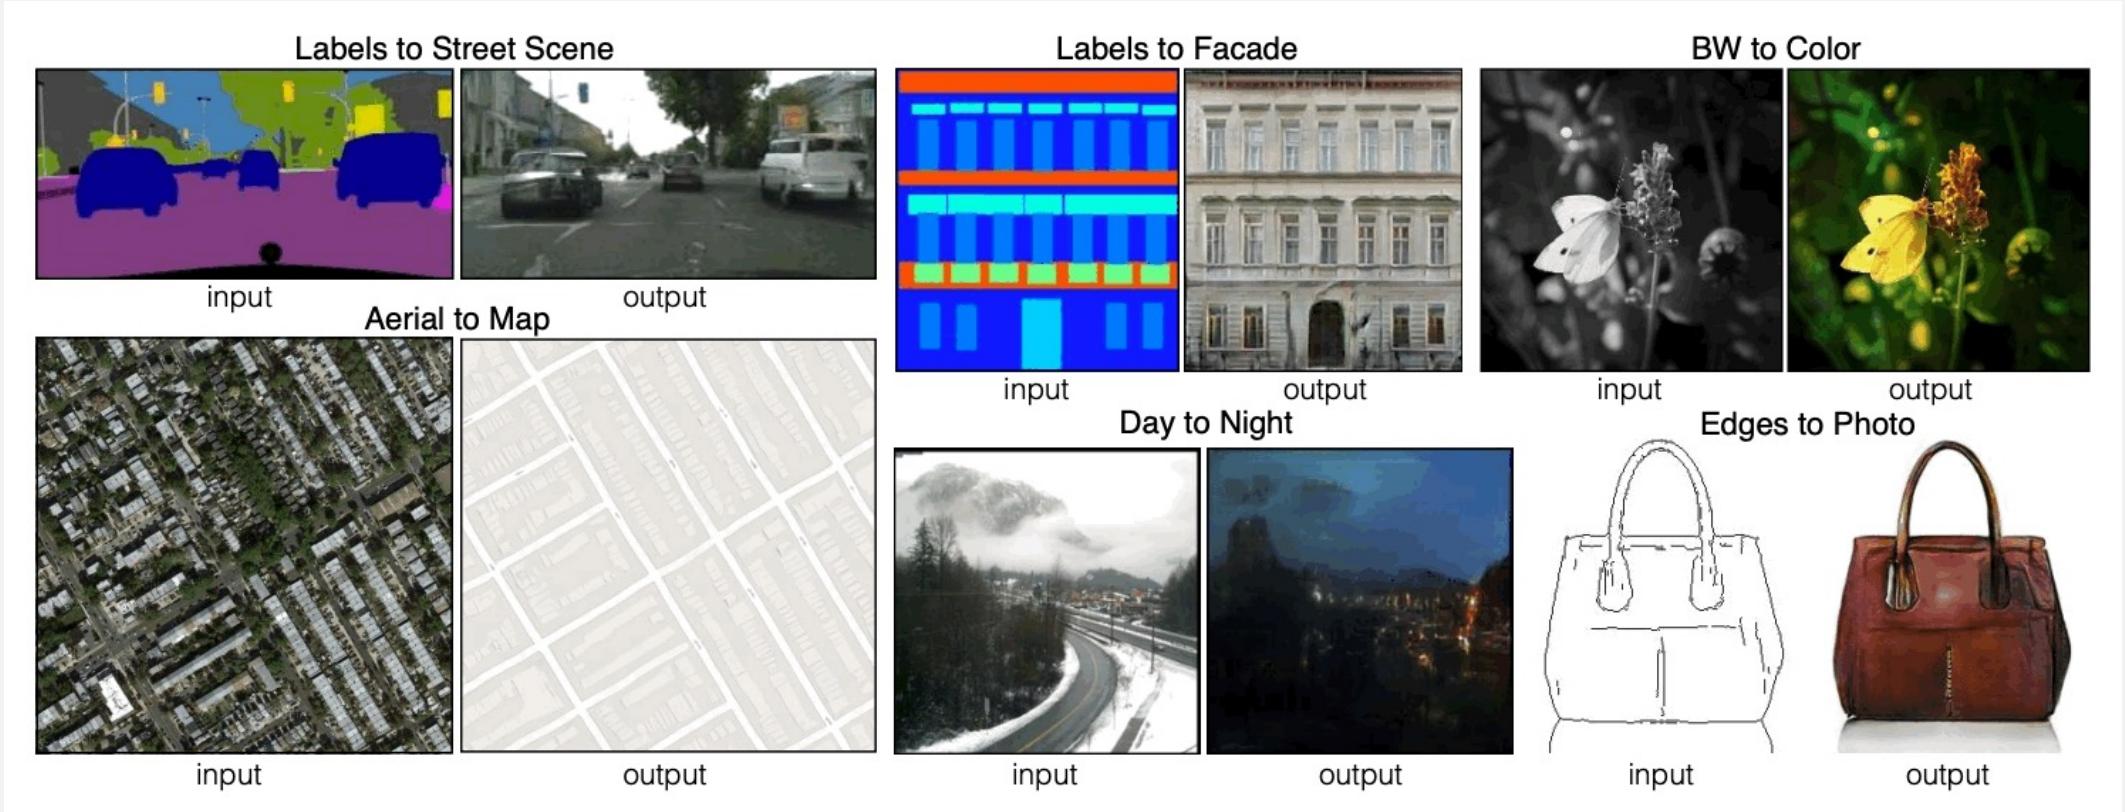


PART II: IMAGE-TO-IMAGE TRANSLATION (PIX2PIX)

- You can find the code in the file:
“ExII Image-to-Image translation
pix2pix.ipynb”

PIX2PIX

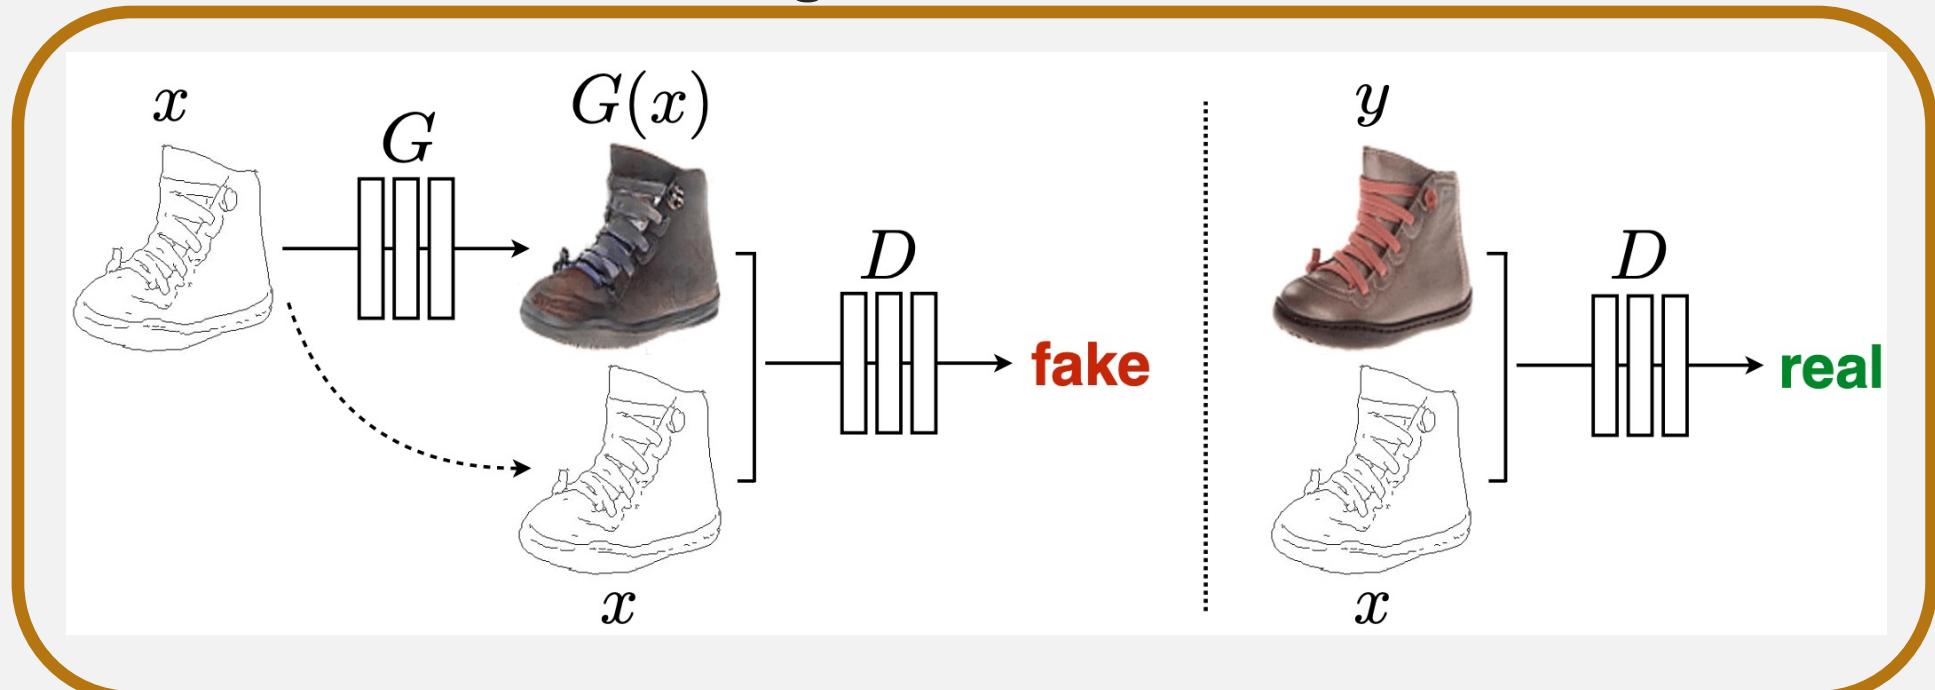
- Image-to-Image Translation with **Conditional Adversarial Networks**
 - Convert an image into another one
 - GAN is conditioned on the input image



PIX2PIX

- A naive approach such as CNN + Euclidean distance as Loss would produce blurry results
- GANs do the following:
 - Set a high level goal: "make output indistinguishable from reality"
 - Automatically learn the loss function to do it

Training Procedure



PIX2PIX

- Objective of a conditional GAN:

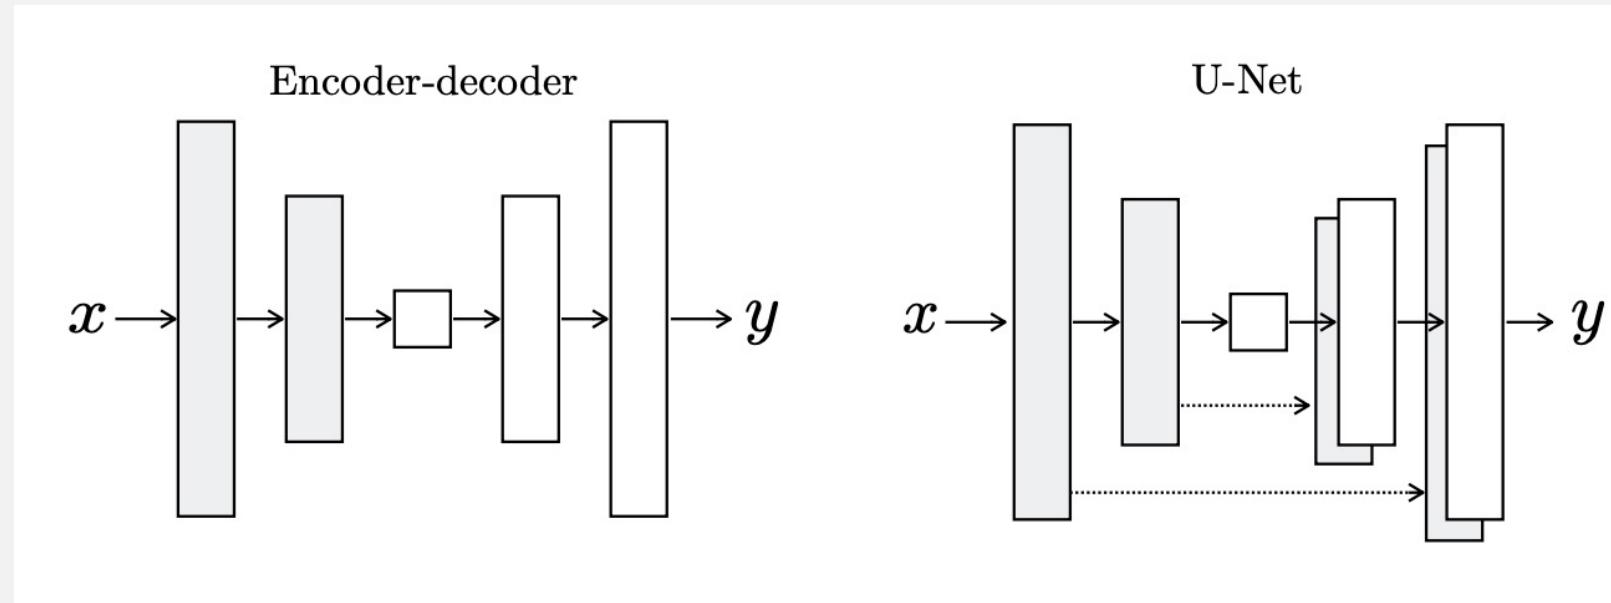
$$\begin{aligned}\mathcal{L}_{cGAN}(G, D) = & \mathbb{E}_{x,y}[\log D(x, y)] + \\ & \mathbb{E}_{x,z}[\log(1 - D(x, G(x, z)))]\end{aligned}$$

- The conditioning is given by the fact that the discriminator observes both input and desired image

PIX2PIX

- **Generator:**

- **U-Net** architecture
- Autoencoder with skipped connections
- both input and output are renderings of the same underlying structure

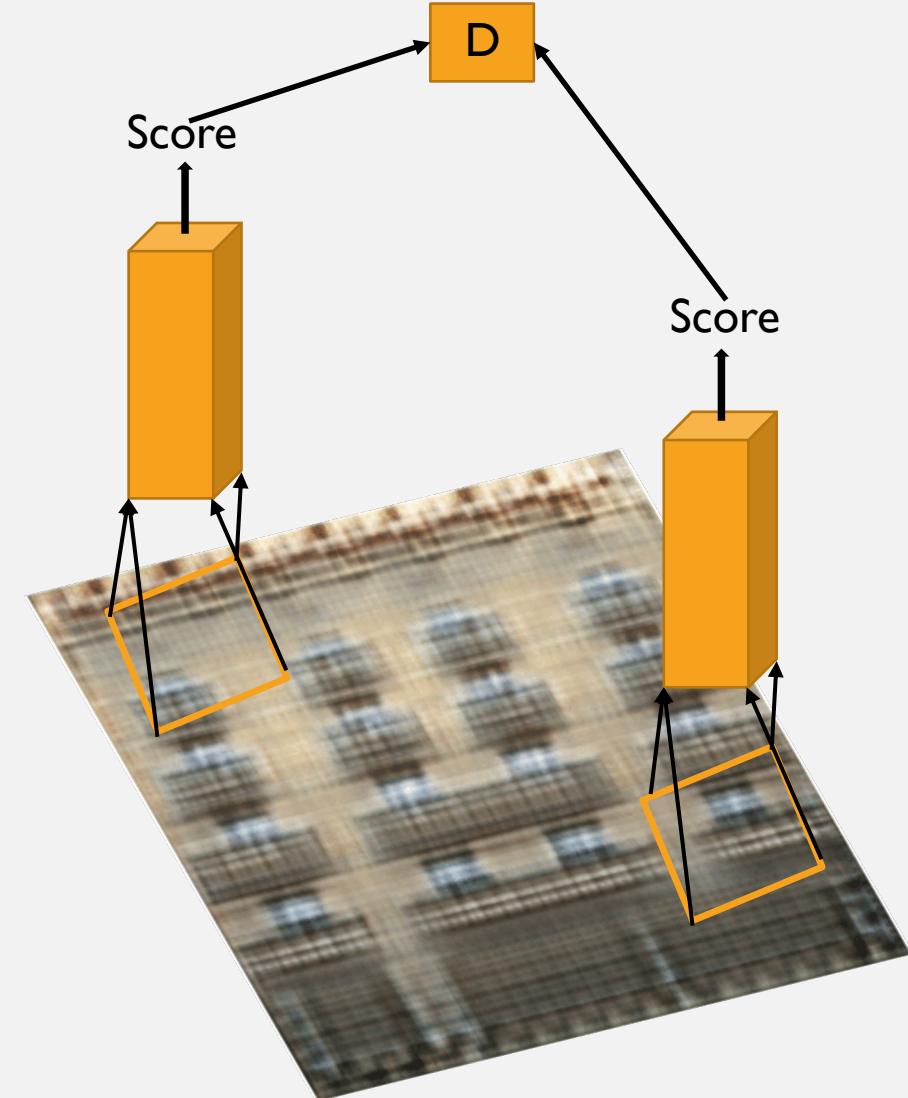


- In many image translation problems a lot of low-level information is shared between input and output
- Skip connections between each layer i and $n - i$ circumvents this issue

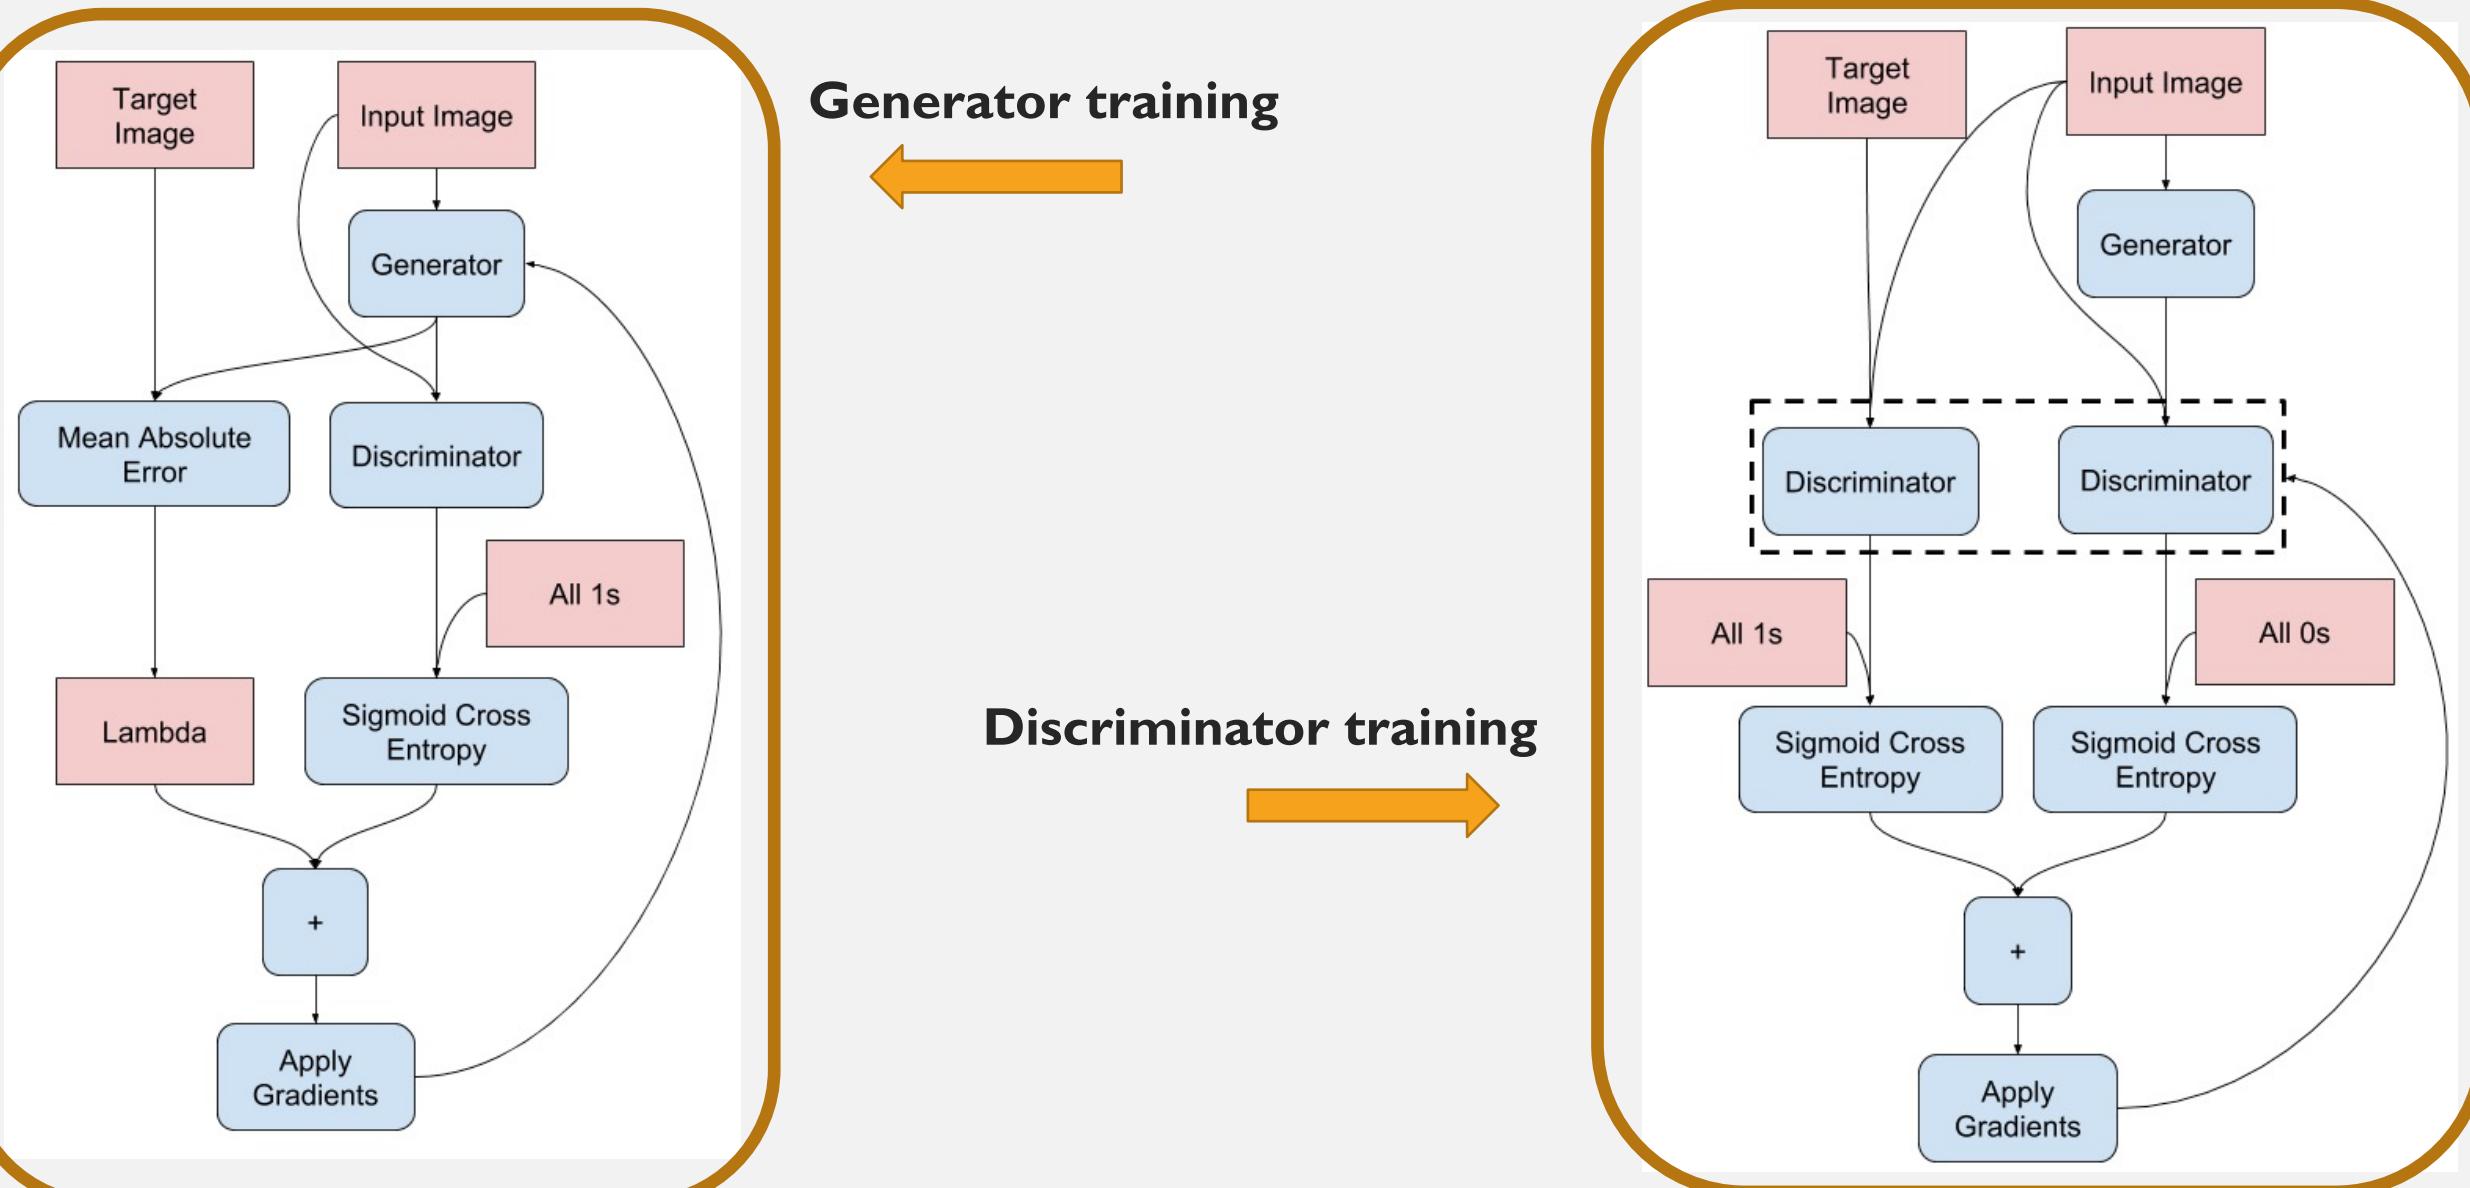
PIX2PIX

- **Discriminator:**

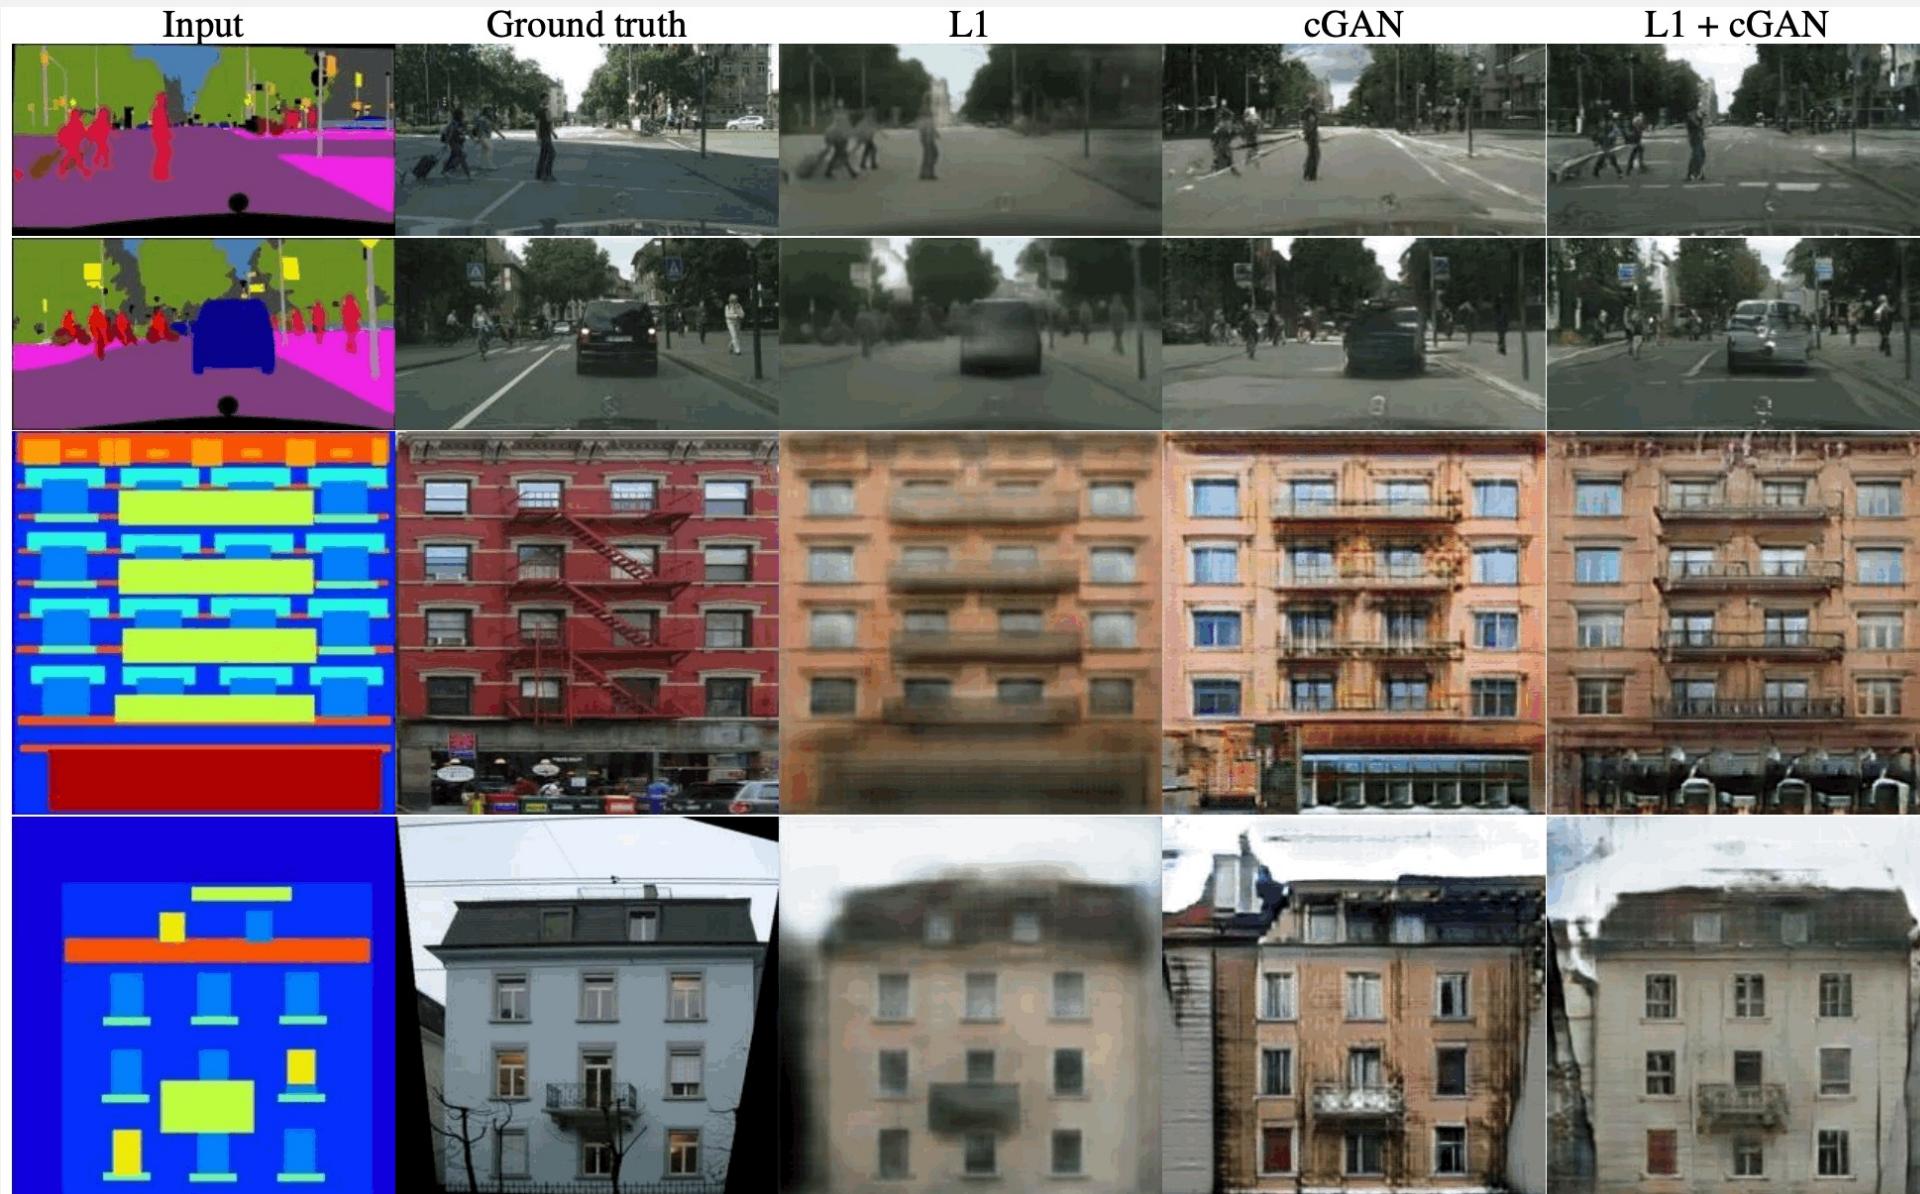
- Patch GAN
 - Penalizes structure at the scale of patches
 - Tries to classify if each $N \times N$ patch in an image is real or fake
 - Run across the image, averaging all responses to provide ultimate output of the Discriminator
 - Can be understood as a form of texture/style loss



PIX2PIX



PIX2PIX



PIX2PIX

Background removal



by Kaihu Chen

Palette generation



by Jack Qiao

Sketch → Portrait



by Mario Klingemann

Sketch → Pokemon



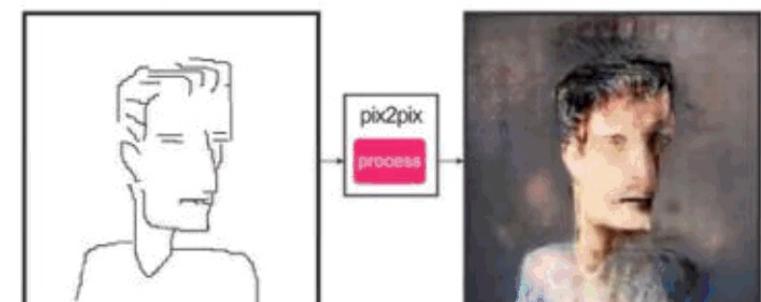
by Bertrand Gondouin

“Do as I do”



by Brannon Dorsey

#fotogenerator



sketch by Yann LeCun

REFERENCES

- Part I
 - Music generation using Recurrent Neural Networks (RNNs)
 - [MIT 6.S191 Introduction to Deep Learning – Lab 1](#)
 - [Music Generation using RNNs \(MIDI\)](#)
 - [Performance RNN](#)
 - [Music Transformer](#)
 - Sketch RNN - [Original paper](#) - [Blog post](#)- [Demo](#)- [Sketch RNN and tensorflow.js](#)
 - LSTM – [blog post](#)

REFERENCES

- **Part II**
 - **Pix2Pix**
 - [Original paper](#) - [Tensorflow tutorial](#) - [Image-to-image Demo](#)
 - GANSynth
 - [Original paper](#) - [Blog post](#) - [Audio Examples](#)